

## **AI/ML Security & Governance — 10 Practical Steps**

1. **Data Minimization** Collect and store only the minimum data necessary for the model. Remove or mask personally identifiable information (PII) before ingestion. This reduces privacy risks and compliance overhead.
2. **Access Control & Audit Logs** Restrict access to sensitive datasets and model endpoints. Maintain immutable audit logs to track who accessed what and when, supporting accountability and investigations.
3. **PII Filtering & Anonymization** Before sending data to external APIs or cloud services, apply PII filters and anonymization. Replace sensitive tokens with pseudonyms to safeguard user identities.
4. **Model Explainability** Document why a model was chosen, its inputs, and its decision-making process. Provide simple explainability notes for business stakeholders to build trust.
5. **Secure Inference** Ensure inference requests are encrypted in transit (HTTPS/TLS). Use private endpoints, authentication, and rate-limiting to prevent misuse of your models.
6. **Bias & Fairness Checks** Continuously test datasets and model outputs for bias. Track fairness metrics and publish mitigation plans as part of release notes to ensure accountability.
7. **Governance Gate for Releases** Introduce a formal release checklist. Each model version must pass governance checks: data quality, bias, privacy, reproducibility, and cost thresholds.
8. **Human-in-the-loop & Escalation** For high-risk use cases (e.g., healthcare, finance), involve human reviewers before final decisions. Define clear escalation workflows for sensitive outputs.
9. **Monitoring & Drift Detection** Deploy monitoring for input distribution shifts and performance decay. Trigger retraining, alerts, or rollbacks when drift or anomalies are detected.
10. **Contracts & Third-party Assessment** When integrating external APIs or third-party models, vet them for compliance and add clear SLA/security clauses. Ensure vendors meet your governance standards.